The Weekend Bookmark

What caught my eye on X this week

24 OCTOBER 2026 · NO. 3 WEEKEND EDITION

INTERPRETABILITY
Untying the neural knot
Page 10

SUPPLY CHAIN
The poisoned package
Page 7

ALTERSVORSORGE
Das Menü fürs Vorsorgedepot
Page 9

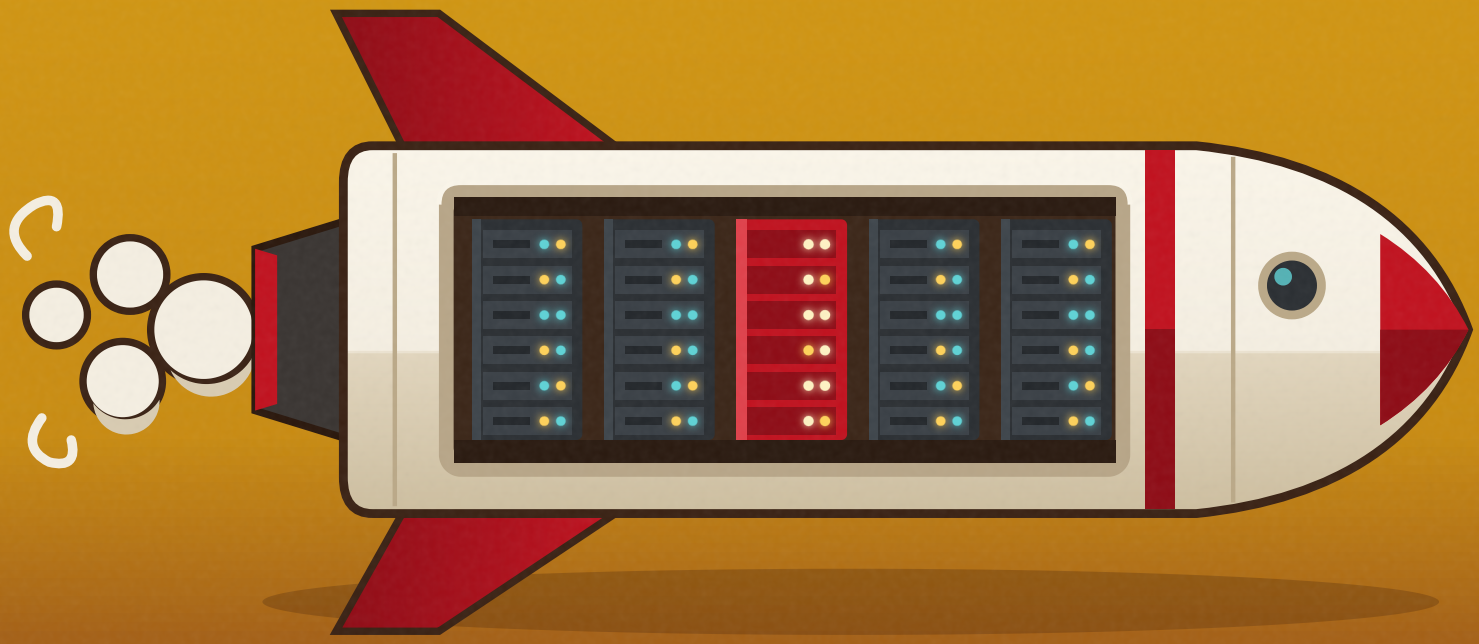

The rocket firm's server farm

ILLUSTRATION: AI-GENERATED FOR THIS ISSUE

ALSO IN THIS ISSUE Leaders 3 | Markets 8 | Faith 13 | The last word 16

Contents

3 LEADERS
Write the memo yourself

4 THE LEAD · COVER
SpaceX's Super Intelligence arm now outsells its space business

5 AI & TOOLS
Anthropic outgrew git, so it built Scion

6 TECH & BUSINESS
A near-record index with 206 battered stocks

7 TECH & BUSINESS
A poisoned LiteLLM release went hunting for AI keys

8 MARKETS
The week in the watchlist

9 POLITICS & SOCIETY
Vom Tech-Plus zum Anleihe-Minus: das Menü fürs Vorsorgedepot

10 IDEAS & CULTURE
The neural knots that may not exist

12 AI & TOOLS
Claude finds the clock behind a famous curve

13 FAITH
Luke's ravens and lilies, and the limits of worry

15 WORTH YOUR TIME
Two long reads

16 THE LAST WORD
The parents' evening has no back row

On the cover: *The rocket firm's server farm*, page 4.

QUOTE OF THE WEEK

"I co-authored early work in mechanistic interpretability at OpenAI and came away pessimistic, because neural networks appeared to fold into arbitrary tangled knots."

@nickcammarata · page 10

PRINTING THIS ISSUE

A4, double-sided, flip on long edge. To saddle-stitch, print as a booklet on four A3 sheets, flip on short edge, fold and staple twice.

The world this week

POLITICS

FROM THE BOOKMARKS

Google Quantum AI demonstrated "magic state cultivation" on physical hardware for the first time, Sundar Pichai (@sundarpichai) wrote, with a 40-fold error reduction on its Willow processor and 99.99 per cent logical fidelity. @sundarpichai

Marius (@gymtrader_) will sein Geld am 19. Oktober wohl von Scalable zur DKB schieben, schrieb er: Neugeld von anderen Banken bringe dort vier Prozent Zinsen; parken wolle er wahrscheinlich 100.000 Euro. @gymtrader_

US equity ETFs drew $1.4 trillion of inflows over the past 12 months, a record, The Kobeissi Letter (@KobeissiLetter) wrote, up from $850 billion in the preceding year; tech funds alone took $65 billion. @KobeissiLetter

Unlimited OCR, a 3B-parameter open-source model from Baidu, reads a 100-page PDF in one continuous pass rather than page by page, @BrooksWhaleX wrote; all figures are the poster's. @BrooksWhaleX

Impeccable (@impeccable_ai) posted a video by @Abdul_wahab_one teasing AI tools: "You asked for better. Your AI added a border." @impeccable_ai

James Cat (@TSLAFanMtl) wrote that, 11 hours into a poll of his followers, the vote was conclusive and he will begin analysing and reporting on SpaceX ($SPCX) alongside Tesla. He said his Tesla coverage would not change and that a merger may soon make the move inevitable anyway. @TSLAFanMtl

Wayne Sutton (@waynesutton) praised "You can just build things", a hardware showcase where ChatGPT users vote on or submit projects. @waynesutton

THE WEEKEND BOOKMARK

Never miss a Saturday.

The week's best bookmarks, edited down to sixteen pages and delivered as a print-ready PDF every Saturday morning.

bkmk.to/subscribe

BUSINESS

FROM THE BOOKMARKS

Varda Space said it was three for three in orbit this October; its W-7 mission carries a protein-degrader cancer drug. Wajahat Mughal (@Wajahat) wrote that microgravity crystallisation promises purer drugs. @Wajahat

Elon Musk (@elonmusk) said Grok's planned Hagglebot for business could improve enterprise cost-efficiency by 5 per cent or more, which he put at up to $6tn a year in savings globally. @elonmusk

Reverse Engineer Anything has crossed 300,000 GitHub stars. Matt Van Horn (@mvanhorn) wrote that it is not about building an open-source Photoshop or "the death of software" but "about unlocking human creativity". @mvanhorn

A PHP flaw lets anyone take control of the computer hosting a website, wrote @sherlock_comms, who said 70 per cent of websites use PHP and the flaw was announced before any patch. @sherlock_comms

Nick Dobos (@NickADobos) wrote that large agent swarms must be prompted "with love"; values, culture, feeling and story-based prompting, he said, is the only way to steer 10 billion decisions a second. @NickADobos

Art Encyclopedia (@artenpedia) posted Salvador Dalí's The Ascension of Christ (1958) on 17 October; it drew 8,817 likes, though the post itself offered nothing beyond the artist, title and year. @artenpedia

ATTENTION

Find the task that matters

Running many agents at once is not the same as progress; the better question is which single task deserves full attention.

PROMPTED BY @trq212

Thariq (@trq212) wrote on Monday 19 October 2026 that most people multitasking with agents "could benefit from slowing down", adding that "most of the stuff you're doing wont really matter".[1] We agree. Running more agents in parallel can feel like productivity; it may be busyness at machine speed. The harder discipline is choosing the task that matters and giving it full attention.[1]

Thariq did not offer a method. The post is brief; it speaks to anyone running several agents at once.[1] His post closes with questions: "what's your P0?" and "what would it look like if you gave it your full attention?"[1] The post does not define P0 or say which tasks qualify, so each reader must decide that for themselves. It also does not say how many agents is too many, nor does it name any tool or product. The advice is narrower than it sounds: slow down, pick the thing that matters, and look at it properly.[1] ■

SOURCES 1 @trq212 bkmk.to/3n2

MODELS

A parts list is not a verdict

Describing how a language model works, as Gregory Kurtzer did in a recent post, does not settle what it can do.

PROMPTED BY @gmkurtzer

Gregory Kurtzer (@gmkurtzer) wrote that a large language model is "a statistical next-token prediction model", an "extraordinarily sophisticated calculator".[1] He is right about the parts. But a parts list settles how a thing works, not what it can do. He concedes as much, giving "full credit" to its capabilities.[1] He notes that its weights stay fixed between inference requests, so it does not learn.[1] He adds that the architectures, mathematics, implementations and research are extensively documented, so anyone can study them.[1] He says he is open to discussing "machine consciousness" in terms of "observable, testable computational properties", and asks how anyone would define it and measure it.[1] Describing the mechanism should start the argument, not end it. ■

SOURCES 1 @gmkurtzer bkmk.to/3n3

WRITING

Write the memo yourself

When colleagues suspect a machine wrote the memo, they stop reading, so firms should make authorship part of the document.

PROMPTED BY @TrungTPhan

ILLUSTRATION

ILLUSTRATION: AI-GENERATED FOR THIS ISSUE

Amazon's six-page memo has a new enemy, according to Trung Phan (@TrungTPhan). Its usefulness, he wrote, "has been nuked by AI-generated text": employees sit in meetings suspecting the prose came from a machine, and tune out.[1] Memos that arouse suspicion earn a tag, AI;DR, which he spells out as "AI, didn't read".[1] His post is an observation, not a study, and it offers no figures on how common the habit is.[1] Phan posted it on Sunday 18 October 2026.[1]

We think the trouble runs deeper than wasted meetings. The point of a long memo is that its author has to think on the page. Full sentences expose a weak argument that a slide can hide, and readers trust a conclusion partly because they can watch someone reach it. That trust is the format's real currency, and it rests on an assumption so obvious it is rarely stated: a person did the thinking. A memo is a promise from its author to its readers: I have thought about this so that you need not. Break that promise, even by appearance, and the reader spends the meeting on doubt rather than on the argument.

Suspicion alone is enough to break it. The readers in Phan's account do not wait for proof that a memo was generated; they simply tune out.[1] That is the cruel arithmetic of doubt. A carefully argued document and a padded, machine-drafted rival now pay the same tax, because a reader cannot tell them apart at a glance. The honest author loses most, and the incentive to bother shrinks with every meeting.

Firms should therefore treat authorship as part of the document. Put the author's name on it, say plainly which passages a model drafted, and expect whoever signs it to defend every paragraph aloud. If the signatory cannot, the memo is not ready, whoever or whatever drafted it. Software can check a draft, tidy a table or catch a contradiction. It should not supply the reasoning the memo exists to display, because once that reasoning is outsourced, the meeting has nothing left to examine.

A memo that colleagues will not read is worse than no memo at all: it wastes the writing and the meeting. If a document deserves six pages of a colleague's attention, it deserves its author's too. Anything less, and AI;DR becomes the most honest review it gets.[1] ■

SOURCES 1 @TrungTPhan bkmk.to/3n1

THE ARGUMENT IN ONE LINE *A memo nobody believes a person wrote is a memo nobody reads; the thinking must be the author's.*

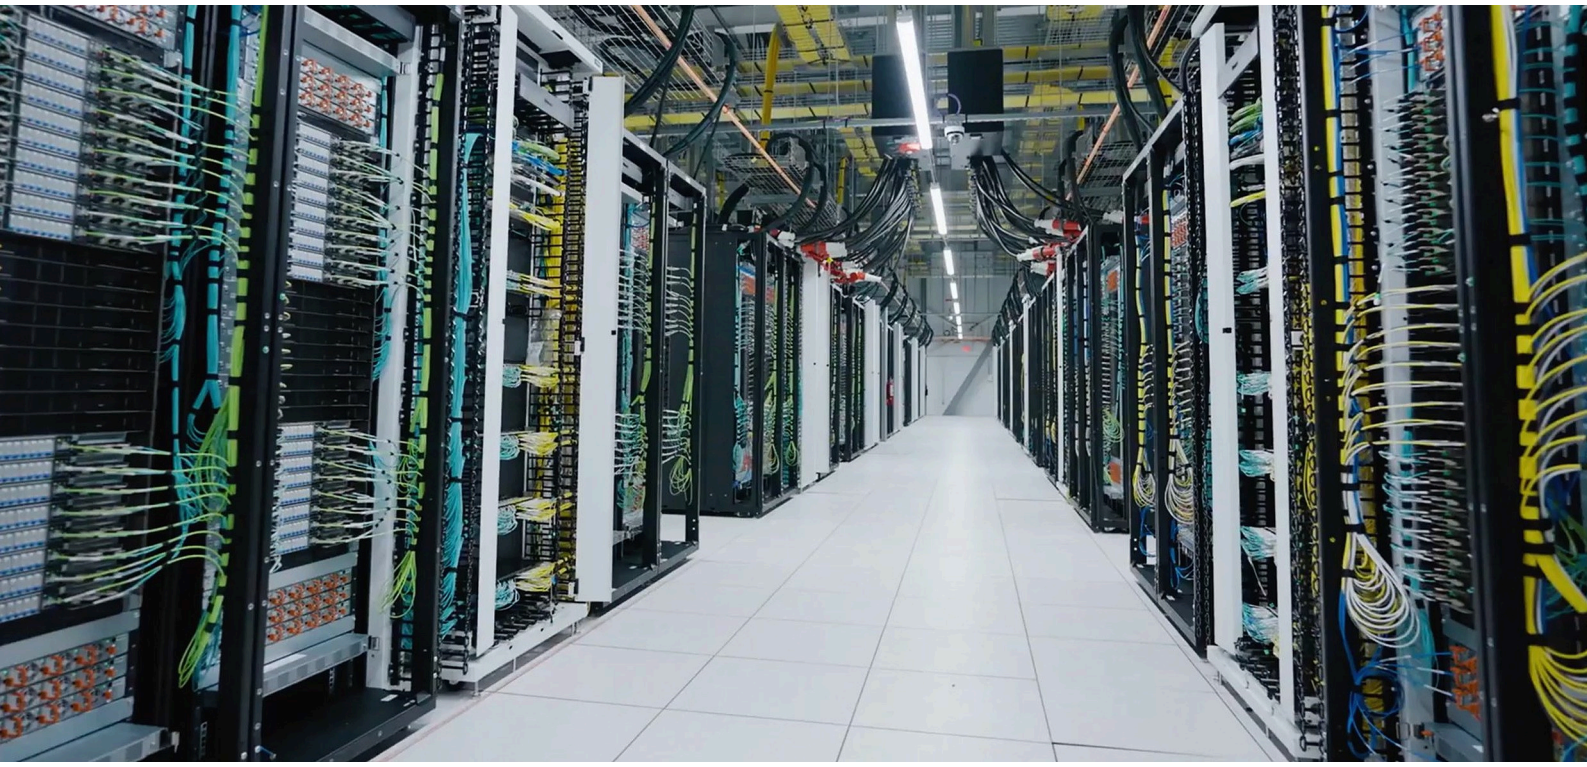

PHOTO: @SAWYERMERRITT

COMPUTE

SpaceX's Super Intelligence arm now outsells its space business

SOURCED VIA @SawyerMerritt

SpaceX's Super Intelligence (SI) segment, formerly labelled AI, now brings in more revenue than its connectivity and space businesses put together. In the third quarter, SI revenue was $7.56B, against $5.25B for Connectivity and $1.77B for Space, Sawyer Merritt (@SawyerMerritt) wrote.[1] SI revenue rose 830 per cent year on year, while group revenue rose 161 per cent to $14.58B.[1] Cursor contributed $608M in revenue between August 14 and September 30.[2]

SI is also profitable at the operating line. It posted operating income of $431M, against a $1.72B loss a year earlier and a $1.26B loss in Q2, Merritt's summary says.[2] amit (@amitisinvesting) went further, writing that "the datacenter business was net profitable"; the summary reports operating income, not net profit, so the stronger claim is his.[3]

> *"The most impressive part of the earnings was that the datacenter business was net profitable"*
>
> AMIT (@AMITISINVESTING)

Capital spending dwarfs profit. Capex reached $19.59B, up 185 per cent, against net income of $4.40B, and $16.44B of the spending, about 84 per cent, went on SI.[2] The company guided to $35B–$40B of capex in Q4 as it accelerates SI compute deployment, and says newly deployed capacity pays back in under 12 months at current pricing.[2] It ended the quarter with about $90B in cash and marketable securities.[1]

The ambition is measured in gigawatts. Nameplate compute capacity reached 1.7 GW, up from 1.4 GW in Q2 and 0.6 GW a year earlier, with nearly 2 GW expected by the end of October and 2.5–3 GW by the end of 2026.[2] For 2027, Merritt's summary says Elon expects to add at least 10 GW, which could take the total to 13 GW or more.[2] Whole Mars Catalog (@wholemars) wrote: "10 gigawatts of additional compute in 2027?!? They have 1.7 gigawatts now..."[4]

Results beat forecasts, though the sources disagree on what those were. Kyle Reidhead (@KyleReidhead) wrote before the results that the Street expected $12.9B of revenue and earnings per share of about $0.11 to $0.16.[5] Merritt gave the revenue estimate as $12.7B.[1] App Economy Insights (@EconomyApp) called it a $1.6B revenue beat, with GAAP earnings per share of $0.35, $0.19 ahead.[6] Reported revenue of $14.58B and EPS of $0.35 exceeded each version.[16]

Next comes the share register. Reidhead wrote that up to 1.3 billion SPCX shares were due to unlock on the Thursday of results week, the largest release yet, worth about $192 billion and enough to grow the tradable float by about 45 per cent.[5] The sources say nothing of how the market will absorb them. SpaceX also guided to Q1 2027 revenue of $33B–$37B.[2] ■

SOURCES 1 @SawyerMerritt bkmk.to/3a1 2 @SawyerMerritt bkmk.to/3a4 3 @amitisinvesting bkmk.to/3a3 4 @wholemars bkmk.to/3a5 5 @KyleReidhead bkmk.to/3a6 6 @EconomyApp bkmk.to/3a2

DEVELOPER TOOLS

Anthropic outgrew git, so it built Scion

ANTHROPIC SAYS GIT'S DESIGN FOR OFFLINE WORK HELD BACK ITS SCALE AS IT MERGED MORE THAN 30,000 CODE CHANGES A DAY, SO A TEAM OF SEVEN BUILT A REPLACEMENT IN FIVE MONTHS.

SOURCED VIA @atmoio

Mo
@atmoio

POST

> In a few years no 100+ person company will use Git anymore. Its design is seriously archaic. Just look at what's happened to GitHub.

13 OCT 2026 bkmk.to/3b1

POST: @ATMOIO

Anthropic has moved its monorepo off git. It wrote on its engineering blog that it built Scion, an in-house version control system, and migrated in five months. In May 2026 it averaged around 2,500 merged code changes a day; by September, more than 30,000. Coding, the post said, is no longer the bottleneck; the developer infrastructure built for human scale is.

Every git clone is a complete repository that works offline, a design the post concedes suits most projects. Anthropic's is not most: its main-branch history is almost 100 GB, with 25 million CI checkouts a day. Scion assumes its clients, nearly all agents and jobs, are always online, so the server handles history, blame and merges; clients fetch only what they need.

Anthropic says Scion's infrastructure footprint is 200x smaller in CPU and memory than its git mirroring setup, and CI alone gained 25 per cent in CPU capacity. A team of seven built it and Antique, a code-review tool.

Mo (@atmoio) went further, writing "In a few years no 100+ person company will use Git anymore."[1] That is his opinion. Anthropic's post is more cautious: it calls its workload unusual and does not expect many organisations to hit these problems soon, though it suspects others using AI coding tools will. ■

SOURCES 1 @atmoio bkmk.to/3b1

CODING AGENTS

A television loses its adverts and gains a dashboard

RYAN DOYLE SAYS CLAUDE CODE TALKED HIM THROUGH HACKING HIS TELEVISION, REMOVING ITS BLOATWARE AND ADDING A SCREENSAVER THAT SHOWS HIS REVENUE AND FLAGS NEW STRIPE PAYMENTS.

SOURCED VIA @dooooyle

Ryan Doyle (@dooooyle) wrote that he "Hacked my TV with Claude Code today", and listed the gains as "No more ads, bloatware, spyware or laggy UI". The method, by his account, was easy. He sent the coding agent the listing for his television, asked whether it could hack it, and it talked him through how.[1]

Doyle added a screensaver that shows his revenue, calendar and weather, along with notifications for new Stripe payments.[1]

In a later post Doyle wrote "This is how fast the TV moves now", crediting Claude with deleting "all the BS laggy bloatware it comes with".[2] The posts do not name the model, explain the steps or say anything about warranties, so readers tempted to copy him have little to go on.[1][2] ■

SOURCES 1 @dooooyle bkmk.to/3b5 2 @dooooyle bkmk.to/3b6

TRY IT THIS WEEKEND

1. Read Anthropic's engineering post on Scion at anthropic.com/engineering/redesigning-source-control. It explains how Scion stores work as changesets in a write-ahead log and computes blame and file history as derived data.
2. If a large git repository is slowing your team down, try the tools git already offers for working with a slice of it: shallow clones, sparse checkouts and promisor remotes. Anthropic says it used all of them, plus internal git mirrors and repository snapshots kept in blob storage, before building Scion.
3. Look up the prior art Anthropic credits. The company says source control for huge monorepos has precedents, notably Google's Piper and Meta's Sapling, and that it built Scion with their lessons in mind.

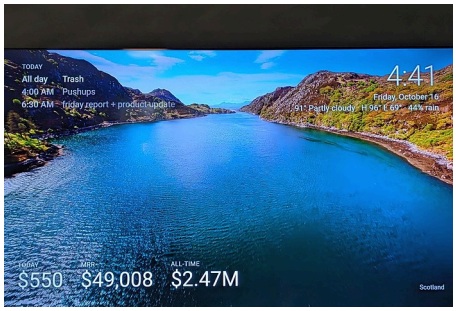

PHOTO: @DOOOYLE

OCR

A 1M-parameter OCR model claims to outrun Tesseract

SOURCED VIA @VikParuchuri

Vik Paruchuri (@VikParuchuri) wrote that Indra, an OCR model with 1M parameters, is 15x faster than Tesseract on CPU and "much more accurate".[1]

On OmniParseBench, Indra lands within 4 points of DeepSeek OCR 2, a 3.4B-parameter model.[2] The figures are the maker's own; neither post cites independent tests.[1][2] ■

SOURCES 1 @VikParuchuri bkmk.to/3b3
2 @VikParuchuri bkmk.to/3b4

AI IN BRIEF

Joe Weisenthal softens on blockchain

Joe Weisenthal (@TheStalwart), who used to mock people calling blockchain's underlying technology interesting, wrote that "maybe they kinda had a point".

@TheStalwart bkmk.to/3k13

MARKETS

A near-record index with 206 battered stocks

ACCORDING TO KOBEISSI, THE S&P 500 IS WITHIN 1 PER CENT OF A RECORD, YET 206 MEMBERS ARE DEEPLY DOWN AND 15 STOCKS HOLD 46.7 PER CENT OF ITS VALUE.

SOURCED VIA @KobeissiLetter

The Kobeissi Letter
@KobeissiLetter

POST

The initial response to this article has been incredible. The S&P 500 is now <1% away from a record high but 206 of its components are down over -30%. Never in history has this happened before. It's the "Bear-Bull Market Paradox." Bookmark this....

16 OCT 2026 bkmk.to/3c3

POST: @KOBEISSILETTER

The S&P 500 is less than 1 per cent from a record high, yet "206 of its components are down over -30%", The Kobeissi Letter (@KobeissiLetter) wrote on October 16th. It says this has never happened and calls it the "Bear-Bull Market Paradox".[1]

On October 21st the newsletter wrote that the 15 largest stocks account for a record 46.7 per cent of the index's market value, against 27.2 per cent during the 2000 dot-com bubble. Nvidia, Apple, Alphabet and Microsoft make up 26.1 per cent, more than the next 11.[2]

The five largest stocks equal the bottom 434 in market cap, which it describes as the largest gap on record. The post does not say how it measured this.[2]

On October 14th the S&P 500 had its worst day since July 29th, and the newsletter said US inflation is rising again, nearly double the Fed's long-term target.[3] On October 15th it wrote the Treasury had failed for months to lower yields until the technology selloff gave relief: "When a drop in equities is the only way to lower yields, you should pay attention."[4] ■

Four firms carry over a quarter of the S&P 500
Share of S&P 500 market value, %, latest

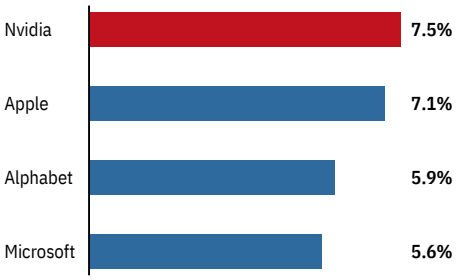

Source: post by @KobeissiLetter

SOURCES 1 @KobeissiLetter bkmk.to/3c3 2 @KobeissiLetter bkmk.to/3c4 3 @KobeissiLetter bkmk.to/3c5 4 @KobeissiLetter bkmk.to/3c2

TRY IT THIS WEEKEND

1. Read The Kobeissi Letter's long article on the "Stealth" bear market, which its posts of October 15th and 16th link to. Treat it as the newsletter's own case: "A Historic Opportunity" in its title is its own framing, not a conclusion the breadth and concentration figures reach by themselves.
2. Do the concentration sum yourself. Add Nvidia's 7.5 per cent, Apple's 7.1 per cent, Alphabet's 5.9 per cent and Microsoft's 5.6 per cent and you reach the 26.1 per cent that Kobeissi cites for the four combined. Then set the top 15 stocks' 46.7 per cent of the S&P 500 beside the 27.2 per cent of the 2000 dot-com bubble, and the remaining 485 stocks' 53.3 per cent.
3. Next time a headline says the S&P 500 is at or near a record, ask the breadth question as well: how many of its members are sharply down? Kobeissi counted 206 components deeply down while the index sat less than 1 per cent from its peak.

PUBLIC DEBT

America's interest bill now tops $1 trillion

THE FEDERAL DEFICIT HIT $1.99 TRILLION IN FISCAL 2026, AND DEBT SERVICE NOW TRAILS ONLY SOCIAL SECURITY AND MEDICARE AMONG WASHINGTON'S EXPENSES.

SOURCED VIA @KobeissiLetter

America's federal deficit widened to $1.99 trillion in fiscal 2026, up 12 per cent from $1.78 trillion last year, according to The Kobeissi Letter (@KobeissiLetter) in a post of 16 October.[1] Spending rose 6 per cent to $7.4 trillion, and net interest passed $1 trillion, making debt service the government's third-largest expense after Social Security and Medicare, the account wrote.[1] It counted the fourth straight year with a deficit around $2 trillion. The account judged that the deficit spending crisis was "out of control".[1]

A later post, dated 17 October, gave a larger figure: interest expense on public debt rose by $144 billion, or 12 per cent, to a record $1.36 trillion.[2] The posts do not say why the figures differ. Both come from the same account, and neither names a government dataset. Interest equalled 25 per cent of total receipts, which rose 3.2 per cent to $5.40 trillion.[2] ■

SOURCES 1 @KobeissiLetter bkmk.to/3c6 2 @KobeissiLetter bkmk.to/3c7

The original supply-chain attack

GIULIO BONASONE, THE TROJANS HAULING THE WOODEN HORSE INTO TROY (1545 (PUBLISHED 1773)). CC0 1.0, THE METROPOLITAN MUSEUM OF ART (OPEN ACCESS)

SUPPLY CHAIN

A poisoned LiteLLM release went hunting for AI keys

MALICIOUS PYPI PACKAGES AND DOCKER IMAGES OF LITELLM 1.107.0 CARRIED A CREDENTIAL STEALER AND A BACKDOOR, LAUNCHED WHENEVER PYTHON STARTS, THE HACKER NEWS REPORTS.

SOURCED VIA @TheHackersNews

Version 1.107.0 of LiteLLM, a Python package, was compromised in a supply chain attack that delivered a credential stealer and a backdoor, The Hacker News reports, citing SafeDep, Socket, StepSecurity and Upwind.[1] It spread through PyPI wheels and Docker images; PyPI has quarantined the package and the malicious images have been removed.[1]

The malware needs no import of LiteLLM: a path file, "cook.pth", runs when Python starts, and its single line, "import cook", launches the collector.[1] It gathers cloud tokens, private keys, AI tool credentials, browser data and environment variables, and can pull saved credentials from LiteLLM proxy deployments.[1]

The data is sent to "litellm[.]dev/healthcheck".[1] StepSecurity said the second payload "forks a background daemon that checks a GitHub dead drop every hour and runs whatever Python the attacker serves."[1] ■

SOURCES 1 @TheHackersNews bkmk.to/3c8

SECURITY

A perfect-ten GitLab flaw reaches back 12 years

SELF-MANAGED SERVERS RUNNING ANY RELEASE SINCE 2014 NEED THE OCTOBER 21 PATCH, INCLUDING THOSE WHOSE ADMINISTRATORS UPDATED BARELY NINE DAYS EARLIER.

SOURCED VIA @TheHackersNews

GitLab released fixes on October 21 for a flaw that lets an attacker run code on a self-managed server without logging in, according to The Hacker News (@TheHackersNews). Tracked as CVE-2026-110264 and rated 10.0 on the CVSS scale, the maximum, it reaches every version back to 7.0, released in 2014. The culprit is Onigmo, Ruby's regular-expression engine, which GitLab ships inside its packages.[1]

Only self-managed installations need to act; GitLab.com is already fixed. GitLab reports no sign of attacks, but has not said what conditions apply and offers no interim workaround. Administrators who patched on October 12 are back in range, because the new fixes sit one patch above that release.[1] ■

SOURCES 1 @TheHackersNews bkmk.to/3c10

BUSINESS IN BRIEF

SpaceX, Tesla and Meta report earnings

The Kobeissi Letter (@KobeissiLetter) listed SpaceX, Tesla and Meta earnings this week, with roughly 15 per cent of the S&P 500 reporting.

@KobeissiLetter bkmk.to/3k15

A prompt priced like a designer

Min Choi (@minchoi) wrote that one Opus 5.5 prompt is "literally a $400/hr motion graphics designer", without saying more.

@minchoi bkmk.to/3k16

Gemini 4 Argon weighs in on elites

@tmuxvim wrote "wow Argon spittin here" of a @PromptLLM post about Gemini 4 Argon on "what separates the normies from the elites".

@tmuxvim bkmk.to/3k17

A Dutch plea to spare a shop

Ernst Lissauer (@ErnstLissauer) wrote, in Dutch, that a shop belonging to "een zuster" should not be plundered, without further context.

@ErnstLissauer bkmk.to/3k18

THE WEEKEND BOOKMARK

Never miss a Saturday.

The week's best bookmarks, edited down to sixteen pages and delivered as a print-ready PDF every Saturday morning.

- Every story credited to its source
- Prints cleanly on any home laser printer
- Booklet mode for folding and stapling
- No feed, no ads, no notifications

bkmk.to/subscribe

MARKETS

The week in the watchlist

29 MARKETS AND STOCKS FROM THIS WEEK'S BOOKMARKS, AND HOW THEY MOVED

Prices as of Thursday 22 October 2026, close

DAX	S&P 500	ZCASH	BITCOIN
25,371.75	7,648.94	€1,089.30	€72,552.99
−0.0%	−1.2%	+0.1%	−1.0%

COMPANY	LAST	WEEK	5 DAYS
Tesla TSLA	$377.53	+4.2%	
Scout24 G24	€85.70	+8.1%	
Apple AAPL	$347.49	+1.6%	
Palantir PLTR	$195.68	−2.4%	
Coinbase COIN	$168.98	−8.1%	
Microsoft MSFT	$520.84	−3.4%	
Alphabet C GOOG	$319.27	−7.4%	
Meta META	$642.86	−12.7%	
Amazon AMZN	$249.17	−2.2%	
Nvidia NVDA	$211.93	−2.4%	
Spotify SPOT	$585.10	+9.7%	
Airbnb ABNB	$175.30	+0.6%	
AMD AMD	$504.25	−10.3%	

COMPANY	LAST	WEEK	5 DAYS
Strategy MSTR	$145.55	−3.8%	
Robinhood HOOD	$96.30	−7.4%	
Cloudflare NET	$307.28	−2.9%	
TSMC TSM	$411.19	−2.6%	
TransDigm TDG	$1,137.86	+0.3%	
Oracle ORCL	$124.37	−7.5%	
The Zcash ETF ZCSH	$32.42	−1.3%	
Siemens Energy ENR	€122.08	−8.7%	
ASML ASML	$1,562.51	−5.7%	
Alphabet A GOOGL	$322.83	−7.6%	
Space Exploration Technologies SPCX	$145.29	−2.2%	
Broadcom AVGO	$322.87	−6.2%	

Chipmakers fell far harder than the market

Share price, 22 Sep 2026 = 100

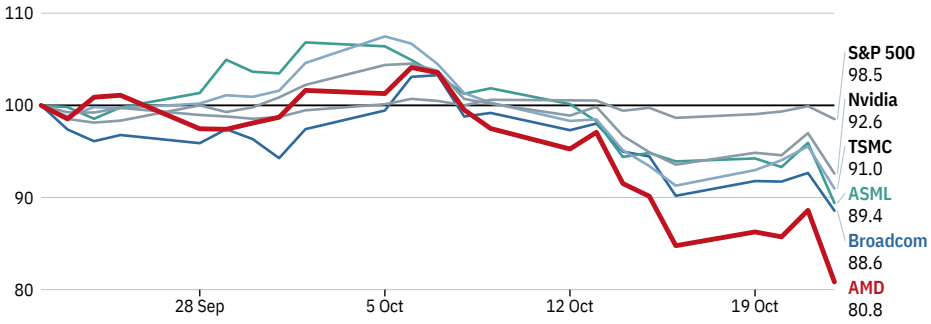

Source: Yahoo Finance

HOW TO READ THIS PAGE

Nineteen of 25 watchlist stocks fell this week

META HAD THE STEEPEST FALL AND SPOTIFY THE BIGGEST RISE, WHILE THE S&P 500 LOST GROUND AND THE DAX ENDED THE WEEK LITTLE CHANGED.

In the week to 22 October, 19 of the 25 stocks on the watchlist fell and six rose, according to Yahoo Finance price data. Spotify led the gainers, up 9.7 per cent to $585.10, followed by Scout24, up 8.1 per cent, and Tesla, up 4.2 per cent.[1]

The falls were steeper. Meta dropped 12.7 per cent to $642.86, the largest fall on the list, ahead of AMD, down 10.3 per cent, Siemens Energy, down 8.7 per cent, and Coinbase, down 8.1 per cent. The data does not say why any of them moved.[1]

Among the benchmarks, the S&P 500 lost 1.2 per cent to close at 7,648.94 points, while the DAX was flat on the week at 25,371.75. Bitcoin fell 1.0 per cent; Zcash rose 0.1 per cent.[1]

MOVERS

Meta −12.7%

Down 12.7% on the week to $642.86.

META

AMD −10.3%

Down 10.3% on the week to $504.25.

AMD

Spotify +9.7%

Up 9.7% on the week to $585.10.

SPOT

Siemens Energy −8.7%

Down 8.7% on the week to €122.08.

ENR

THE WEEK AHEAD — RESULTS DATES

Tue 27	Robinhood results
Wed 28	Alphabet C results
Wed 28	Alphabet A results
Wed 28	Microsoft results
Thu 29	Amazon results
Thu 29	Coinbase results

SOURCES 1 Prices: Yahoo Finance chart bkmk.to/3m1

Prices: Yahoo Finance chart API (29); earnings: Nasdaq; closing prices to 22 October 2026. * Results date estimated. Not investment advice.

ALTERSVORSORGE

Vom Tech-Plus zum Anleihe-Minus: das Menü fürs Vorsorgedepot

@Techaktien1 listet die Fonds für das Altersvorsorgedepot 2027 bei Trade Republic auf; ihre Fünfjahreszahlen reichen von +158 Prozent bis -14 Prozent und sind keine Prognose.

SOURCED VIA @Techaktien1

Die Spitze gehört der Technologie. Laut @Techaktien1 (Techaktien) legte Artificial Intelligence & Big Data über fünf Jahre um 158 Prozent zu, MSCI World Information Technology um 144 Prozent.[1] Der MSCI World kommt auf 77 Prozent, der DAX auf 58 Prozent.[1] Stichtag und Berechnungsmethode nennt der Beitrag nicht.[1]

"Anders als der erstplatzierte Artificial Intelligence & Big Data ETF (0,35% TER) enthält er keine chin. Werte."

@TECHAKTIEN1

Seine eigene Wahl begründet der Autor offen. Er werde "...zu Beginn des Jahres wahrscheinlich für 1.800€ den Xtrackers MSCI World Information Technology A113FM holen", schrieb @Techaktien1. Der Fonds mit 0,25 Prozent TER ähnele dem Nasdaq 100, enthalte anders als der Spitzenreiter aber keine chinesischen Werte.[1] Der Nasdaq 100 (plus 107 Prozent) steht nicht zur Auswahl, schrieb der Autor.[1]

Vergangene Renditen sind keine Prognose: Die Liste zeigt, was in fünf Jahren geschah, nicht, was bis zur Rente geschieht.[1] Die Bookmark gibt keine Anlageempfehlung. ■

SOURCES 1 @Techaktien1 bkmk.to/3d1

SCREENSHOT: WELT, VIA @TECHAKTIEN1

Technologie liegt über fünf Jahre vorn

Fünfjahresrendite einer Auswahl der Fonds im Altersvorsorgedepot von Trade Republic, %

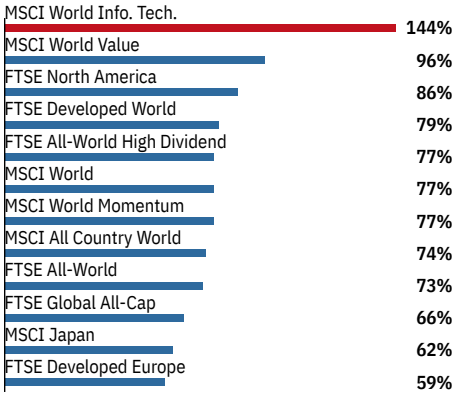

Quelle: Post von @Techaktien1

AM WOCHENENDE PRÜFEN

1. Die Fünfjahreszahlen von @Techaktien1 neben die eigene Übersicht legen und breite Indizes wie MSCI World (+77 Prozent) und FTSE All-World (+73 Prozent) mit Spezialthemen wie Artificial Intelligence & Big Data (+158 Prozent) vergleichen.
2. Die Kosten der engeren Wahl vergleichen: @Techaktien1 nennt für den Xtrackers MSCI World Information Technology 0,25 Prozent TER und für den Artificial Intelligence & Big Data ETF 0,35 Prozent TER.
3. Prüfen, was in den Fonds steckt: Laut @Techaktien1 enthält der Artificial Intelligence & Big Data ETF chinesische Werte, der Xtrackers MSCI World Information Technology nicht.

VERBRAUCHER

Klarna und PayPal sollen Ratenkäufer auf Bonität prüfen

Belegt ist die Änderung zum 20. November bisher nur durch einen scherzhaften Post, der vorrechnet, wie man den Tag davor auf Pump ausreizt.

SOURCED VIA @ALG2Horse

Wer bei Klarna oder PayPal in Raten kauft, soll ab dem 20. November auf seine Bonität geprüft werden. So schrieb es @ALG2Horse (Horse Trading) am 17. Oktober und nannte $KLAR und $PYPL.[1] Weitere Belege und Unternehmensstellungnahmen fehlen.

Weil GTA 6 am 19. erscheine, schrieb er: "ich hab also genau einen Tag, an dem ich noch ohne Prüfung auf Pump zocken darf."[1] Hinzu kommen ein zweiter Controller, PS Plus und Energydrinks.[1] Er landet bei knapp 340 € in vier Klarna-Verträgen, bezahlt über seine Trade Republic Mirror Karte mit 1 Prozent Saveback. Macht 3,40 €; in einen ETF angelegt, bei 7 Prozent nach 30 Jahren 25,88 €. "damit ist ein Teil des Spiels bereits bezahlt."[1] ■

SOURCES 1 @ALG2Horse bkmk.to/3d2

POLITICS IN BRIEF

An essay blames the state for high prices

Brivael Le Pogam (@brivael) wrote that young protesters blame capitalism for costs created by the state, citing Mark Perry's chart of US prices since 2000. Hospital services, tuition, childcare and housing rose far above inflation, he said, while TVs, software and phones fell; he put French public spending at 57 per cent of GDP.

@brivael bkmk.to/3k21

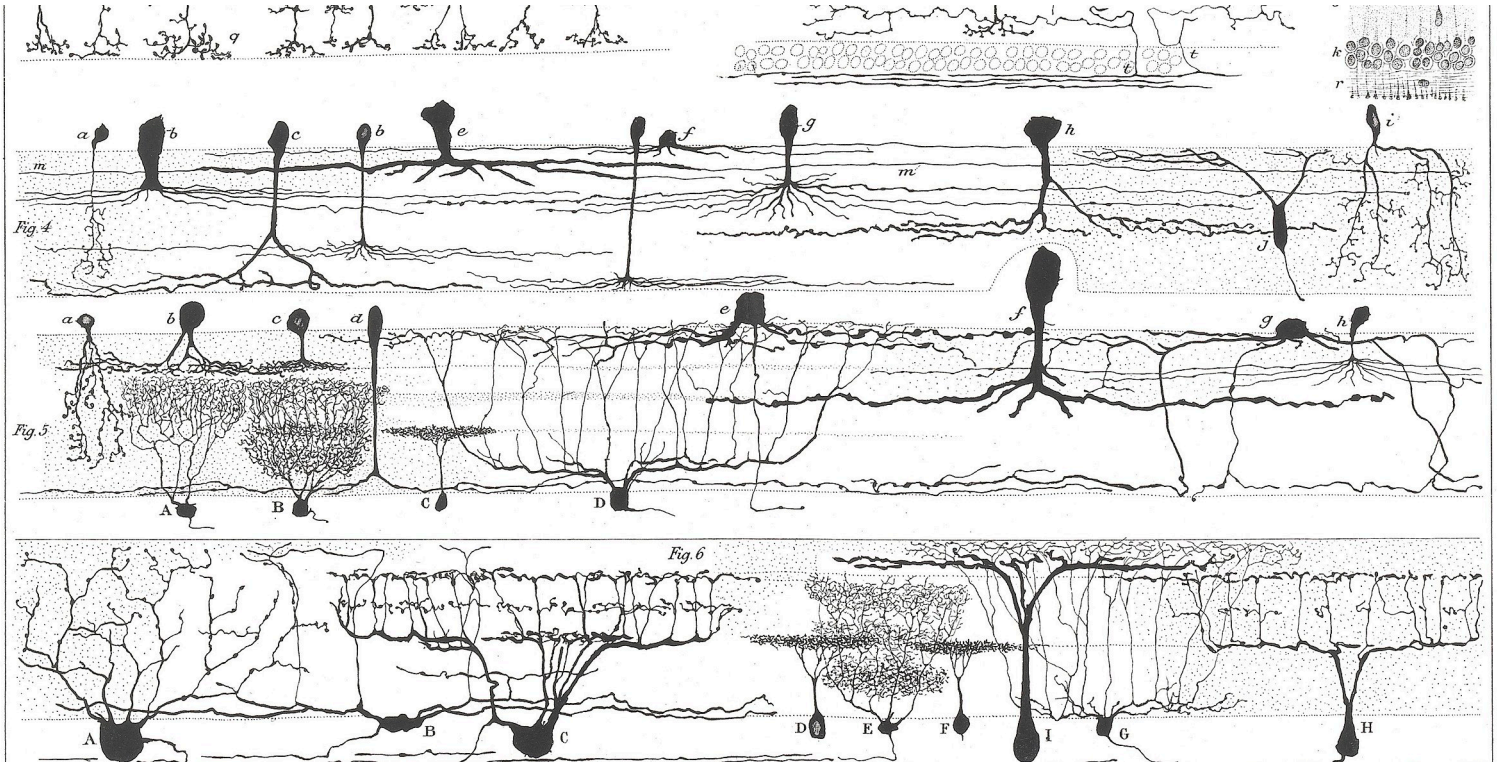

WELLCOME COLLECTION, LINE DRAWING OF THE RETINA. CC BY 4.0, VIA WIKIMEDIA COMMONS

INTERPRETABILITY

The neural knots that may not exist

AN EARLY PESSIMIST ABOUT UNDERSTANDING NEURAL NETWORKS NOW ARGUES, WITH FRESH ESTIMATES FROM VISION AND LANGUAGE MODELS, THAT THEIR NEURONS ARE MOSTLY LEGIBLE AND THAT THIS MATTERS FOR ALIGNMENT.

SOURCED VIA @nickcammarata

Nick Cammarata (@nickcammarata) co-authored early work in mechanistic interpretability at OpenAI and came away pessimistic, because neural networks "appeared to fold into arbitrary tangled knots", he wrote.[1] Now, he added, with modern agentic tools "the picture looks totally different": the vast majority of neurons in models of up to trillions of parameters are legible, and their circuits "beautiful".[1]

The field's leading explanation for polysemantic neurons was superposition: a network that runs out of neurons folds unrelated things together because they do not interfere. The famous example responded to both cat faces and car fronts. The remedy was a second model, such as a sparse autoencoder, to unfold the network.

> *"with modern agentic tools and a much closer look, the picture looks totally different"*
>
> @NICKCAMMARATA

In Neural Networks Are Already Clean, Cammarata argues the canonical evidence was misread. He blames architectural problems and weak tools for InceptionV1's incoherent late layers; with ordinary fixes and no change to model size or training data, he writes, 99–100 per cent of neurons in its last layer are easily interpretable. The 2019 car-and-dog circuit he and Chris Olah found he now reads as "ordinary generalization and composition".

In preliminary audits of Qwen3-32B, he and colleagues estimate that roughly 90 per cent of neuron poles have one main job, about 7 per cent two related jobs and 2–4 per cent are not yet understood.

The stakes are alignment. He wrote that he does not think alignment can be solved without knowing what is being built.[1] He now expects ambitious interpretability to be solved by the time it is most needed, possibly soon: "the simple thing, at absurd scale, works".[1] The claims rest on preliminary audits. If they hold, the knots were mostly in the tools. ■

SOURCES 1 @nickcammarata bkmk.to/3e1

NEUROSCIENCE

The brain flickers between seen and unseen

A preprint shared by Doris Tsao reports that visual and prefrontal regions switch rapidly between the image perceived and the one suppressed from awareness.

SOURCED VIA @doristsao

A preprint shared by Doris Tsao (@doristsao), and led by @HesseJanis, reports that the brain does not hold steady on the image it perceives. In face patches of inferotemporal and prefrontal cortex, activity alternated about 15 times per second between conscious and suppressed micro-epochs, even while the stimulus and the percept stayed constant.[1]

The team used Neuropixels recordings during binocular flash suppression, in which one of two images is consciously perceived while the other is suppressed from awareness. In conscious micro-epochs the authors found a pure signal of conscious perception, clear enough to reconstruct the perceived image from neural activity. In suppressed micro-epochs the signal flipped to track the hidden image, which could also be reconstructed.[1]

> *"brings us significantly closer to the recipe for generating conscious awareness"*
>
> DORIS TSAO (@DORISTSAO)

The epochs looked symmetrical, yet they could be told apart, with feature tuning amplified during conscious ones. Previous studies had found no pure signal of conscious perception anywhere in the brain. Prefrontal recordings gave qualitatively similar results. The authors conclude that conscious access "requires temporal selection beyond localization".[1]

The bigger claims are Tsao's own. She wrote that the work "overturns decades of thinking" on the biological basis of consciousness and "brings us significantly closer to the recipe for generating conscious awareness".[1] The abstract is more modest: inferotemporal and prefrontal cortex interleave conscious and suppressed content. It is a preprint, and the authors declare no competing interest. ■

SOURCES 1 @doristsao bkmk.to/3e3

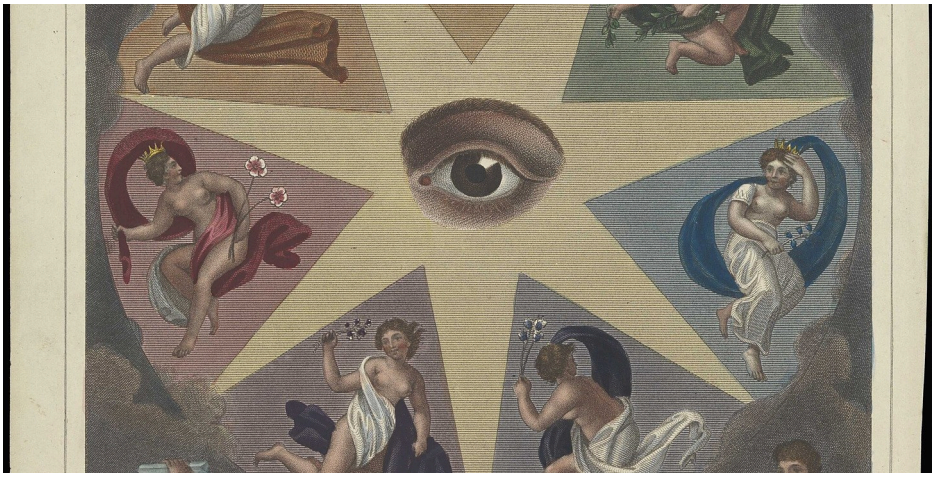

Optics, an allegorical print of the eye (1820)

A.D.M; JOHN CHAPMAN, OPTICS RMG F8609 (1820). PUBLIC DOMAIN, VIA WIKIMEDIA COMMONS

COFFEE

Most decaf is still made with a solvent

A post by @TipOfTheBanana explains how the solvent works, what a 2020 lab test found and where the FDA stands on a petition.

SOURCED VIA @TipOfTheBanana

Most decaf coffee is made with methylene chloride, a solvent that used to be in paint stripper, wrote @TipOfTheBanana.[1] It extracts caffeine and little else.[1] Green beans are steamed, washed in it and steamed again before roasting; the industry calls this the European method.[1] The FDA allows up to 10 parts per million.[1] A 2020 lab test found the solvent in 10 of 24 decaf brands, at about 0.05 to 0.1 ppm, around 100 times below that limit, the post said.[1]

The EPA banned most uses of methylene chloride in 2024, but food falls to the FDA, where a petition to remove it has sat since December 2023.[1] Comments reopened in May; the FDA has not decided.[1] Three alternatives exist: ethyl acetate, pressurised carbon dioxide and the solvent-free Swiss Water process.[1] Peet's says all of its decaf is water process; Starbucks has told reporters most of its decaf is the direct contact method, meaning solvent.[1] ■

SOURCES 1 @TipOfTheBanana bkmk.to/3e5

TRY IT THIS WEEKEND

1. Read the bag: @TipOfTheBanana wrote that Swiss Water or water process on the label means no solvent touched the beans, while a bag that says nothing leaves you not knowing.
2. Look past trade names: the post said ethyl acetate, which the trade calls the natural method, is a solvent that can be made from sugarcane, and the European method uses methylene chloride.
3. Compare brands: according to the post, Peet's says all of its decaf is water process, while Starbucks has told reporters most of its decaf is the direct contact method, meaning solvent.
4. Know the figures: the FDA allows up to 10 parts per million of methylene chloride, and the 2020 test found it in 10 of 24 brands at about 0.05 to 0.1 ppm.

FROM THE REPLIES

A supercut of Tesla's Roadster launch

@stevenmarkryan posted a launch-event supercut headed "Tesla Roadster LITERALLY FLIES", while promoting his Patreon and merchandise; the post gives no specifications.

@stevenmarkryan bkmk.to/3k20

CRYPTOGRAPHY

Claude finds the clock behind a famous curve

For about 25 years the NSA's unexplained seeds for P-256 and its sibling curves fed talk of backdoors.

SOURCED VIA @jarrodwatts

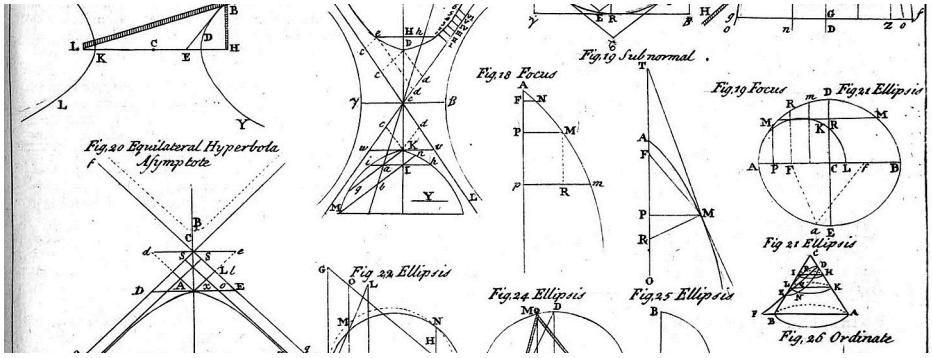

Conic sections, not elliptic curves: a plate from the 1728 Cyclopaedia

TABLE OF CONICS, CYCLOPAEDIA, VOLUME 1, P 304, 1728. PUBLIC DOMAIN, VIA WIKIMEDIA COMMONS

Anthropic says Claude has solved a 25-year-old puzzle. In under half an hour, Claude found that the seeds behind the NIST curves P-192, P-224, P-256, P-384 and P-521 came from the C library's lrand48 random-number generator, started from a Unix timestamp.

NIST adopted ECDSA in 2000 in FIPS 186-2, but the NSA supplied the curves' seeds without explanation, leaving people wondering about deliberate backdoors.

Jarrod Watts (@jarrodwatts) wrote that the main theory for Satoshi's choice of secp256k1 for Bitcoin was distrust of P-256, which Apple's Secure Enclave supports.[1]

The best earlier story, from NSA employee Jerry Solinas, was that the seeds were hashes of a humorous passphrase. That proved wrong. Claude spotted a pattern in the seeds' top bits, recovered the generator's state and traced the prime-curve seeds to timestamps from May 1997 to April 1998.

Anthropic says this is no break, attack or weakness, and that it is more confident the curves carry no backdoors. Watts wrote that this does not yet fully prove there is no backdoor.[1] ■

SOURCES 1 @jarrodwatts bkmk.to/3b7

TRY IT THIS WEEKEND

1. Install the skill with the Skills CLI by running npx skills add https://github.com/cloudflare/security-audit-skill --skill security-audit, adding --global for a user-level installation. Run npx skills --help for agent-selection and non-interactive options. You will also need Node.js for the zero-dependency validators that check the findings and coverage-ledger files, and a coding agent whose model supports tool use and parallel sub-agents.
2. A direct audit request triggers full audit mode. Unless you name a directory, reports go to ~/security-audit-skill/<repo-name>/run-<N>, and the workflow writes inside the target repository only when you choose a directory that version control ignores.
3. Set up an OS-enforced sandbox before letting the skill build or test target code. The sandbox must disable external networking, enforce resource limits and allow writes only to assigned scratch paths. Without it, leads stay marked as needs_validation. Then read REPORT.md and NEEDS-VALIDATION.md and run the audit again: in Cloudflare's tests, a single run found roughly half of what repeated runs found.

SECURITY TOOLS

Cloudflare lends coding agents a sceptical security auditor

An MIT-licensed skill runs coding agents through six phases of hunting and disproof, and @jurbed urges running it with free models too.

SOURCED VIA @jurbed

Cloudflare has published a coding-agent skill that turns an agent into a security auditor across six phases, from reconnaissance to reporting. Its defining habit is organised doubt: each candidate flaw goes to a fresh verifier that tries to disprove it, and confirmed, needs_validation and rejected records are written to findings.json. Reconnaissance maps architecture, trust boundaries and input surfaces; isolated hunters work through units in a coverage ledger, with critics looking for gaps.[1] The verdicts are distinct: a confirmed record needs a complete source trace, while needs_validation names the exact unresolved fact and carries no severity.[1] Reports are derived as REPORT.md, FINDINGS-DETAIL.md and NEEDS-VALIDATION.md.[1]

> *"Run it on all your software. Run it with the free models. Use it even if you used a model to audit something"*
>
> @JURBED

Juraj Bednar (@jurbed) went further than the repository does. He wrote: "Works in all agent harnesses and is model independent." He urged running it with free models, even on code a model has already audited, because "it will find much more things".[1] The repository is more cautious: it asks for a model that supports tool use and parallel sub-agents, plus an OS-enforced sandbox. Without one, it declines to execute target code. ■

SOURCES 1 @jurbed bkmk.to/3b9

JOHN JAMES AUDUBON, RAVEN, PLATE 101 OF THE BIRDS OF AMERICA (1827–38). PUBLIC DOMAIN, VIA WIKIMEDIA COMMONS

WEEKEND ESSAY

Luke's ravens and lilies, and the limits of worry

A passage from Luke 12 sets birds and lilies against the household worries of food and clothing, and ends by redirecting the reader's attention to God's Kingdom.

HELENA DE KAY, WATER LILIES (CA. 1870–77). CC0 1.0, THE METROPOLITAN MUSEUM OF ART (OPEN ACCESS)

THE PASSAGE

Luke 12:22-31 (WEB). The text, with commentary, is overleaf on page 14.

Luke 12:22 opens with an instruction to disciples: "don't be anxious for your life, what you will eat, nor yet for your body, what you will wear." Two worries are named, food and clothing, and a reason follows at once: "Life is more than food, and the body is more than clothing." The argument moves from the ordinary to the larger, and the rest of the passage keeps that order, passing from meals and garments to the question of what a life is organised around.[1]

The first image is the raven. In the World English Bible's wording, ravens "don't sow, they don't reap, they have no warehouse or barn, and God feeds them." The vocabulary is that of a household economy: planting, harvest, storage. Verse 24 then draws a comparison of worth: "How much more valuable are you than birds!" The text does not say here that people should stop sowing or storing. It argues from value, and what that implies for work is a matter of interpretation, on which readers differ.[1]

Verses 25 and 26 change method, from comparison to a practical test. "Which of you by being anxious can add a cubit to his height?" the passage asks, and follows with: "If then you aren't able to do even the least things, why are you anxious about the rest?" The questions ask whether worry produces anything, not whether it is shameful. Verse 28 nonetheless adds a note of reproach, "O you of little faith", and the two registers sit side by side.[1]

The lilies carry the passage's most memorable line: "Consider the lilies, how they grow." They "don't toil, neither do they spin", yet Solomon "in all his glory was not arrayed like one of these." Verse 28 sharpens the image by turning to the grass, "which today exists, and tomorrow is cast into the oven." The promise is therefore not that beauty lasts. The text claims that God clothes even short-lived grass, and from that it argues that the people addressed will be clothed too.[1]

The passage ends with a contrast and an instruction. "For the nations of the world seek after all of these things, but your Father knows that you need these things." Then: "But seek God's Kingdom, and all these things will be added to you." Verse 30 says the Father knows what is needed; verse 31 says these things will be added. How literally to read that promise, as material assurance or as a reordering of priorities, is something interpreters weigh differently, and the passage itself does not settle it.[1] ■

ALSO IN FAITH

A KGB defector's four stages, revisited

Krzysztof Szczawinski (@Kristof_Poland) revisited Yuri Bezmenov, the former KGB propagandist whose 1984 interview described subverting a country in four stages: demoralisation, destabilisation, crisis and normalisation. He wrote that not every number Bezmenov gave can be verified, but that the mechanism holds and has outlived the Soviet Union.

@Kristof_Poland bkmk.to/3k23

A QUESTION TO TAKE AWAY

What would it mean to hold your plans loosely without ceasing to care about the work in front of you?

SOURCES 1 Luke 12:22-31 (WEB) bkmk.to/3f1

LUKE 12:22-31

22 He said to his disciples, “Therefore I tell you, don’t be anxious for
your life, what you will eat, nor yet for your body, what you will wear.
23 Life is more than food, and the body is more than clothing.
24 Consider the ravens: they don’t sow, they don’t reap, they have no
warehouse or barn, and God feeds them. How much more valuable
are you than birds! 25 Which of you by being anxious can add a cubit
to his height? 26 If then you aren’t able to do even the least things,
why are you anxious about the rest? 27 Consider the lilies, how they
grow. They don’t toil, neither do they spin; yet I tell you, even
Solomon in all his glory was not arrayed like one of these. 28 But if
this is how God clothes the grass in the field, which today exists, and
tomorrow is cast into the oven, how much more will he clothe you, O
you of little faith? 29 Don’t seek what you will eat or what you will
drink; neither be anxious. 30 For the nations of the world seek after
all of these things, but your Father knows that you need these things.
31 But seek God’s Kingdom, and all these things will be added to you.

WORLD ENGLISH BIBLE · PUBLIC DOMAIN

COMMENTARY

The passage argues rather than merely consoles. Addressing his disciples, the speaker says “Life is more than food, and the body is more than clothing” (Luke 12:23). He then points to ravens that “don’t sow, they don’t reap,” and to lilies that “don’t toil, neither do they spin,” yet outdo Solomon “in all his glory” (Luke 12:24, 27). The reasoning runs from small to large: “How much more valuable are you than birds!” (Luke 12:24).[1]

Worry gets a practical rebuttal: “Which of you by being anxious can add a cubit to his height?” (Luke 12:25). The text does not call effort pointless; it redirects it. Verse 30 contrasts “the nations of the world” chasing necessities with a Father who knows them, and verse 31 gives the alternative: “But seek God’s Kingdom, and all these things will be added to you.” How literally to apply that is for readers to judge.[1] ■

SOURCES 1 Luke 12:22-31 (WEB) bkmk.to/3f1

HISTORY

What a Galilean household list reveals in Luke

Luke 12:22-31, in the World English Bible, lists the work and goods of a household economy.

Luke 12:22-31 (WEB) names food, clothing, sowing, reaping, a warehouse, a barn, spinning and an oven. It does not date these practices or place them in a particular village. What it offers is a catalogue of the things its speaker treats as the stuff of daily worry.[1]

The first list concerns grain. The text says of the ravens: “they don’t sow, they don’t reap, they have no warehouse or barn, and God feeds them.” The verbs run from planting through harvest to storage, and the birds stand outside every stage. That tells us how the passage pictures security, namely a full barn. It tells us nothing about any real farm or harvest.[1]

The second list concerns cloth. The text says of the lilies: “They don’t toil, neither do they spin; yet I tell you, even Solomon in all his glory was not arrayed like one of these.” Spinning stands here for the labour behind clothing, and Solomon for the highest example of dress. The passage does not say which flower it means. “Lilies” is the WEB’s word.[1]

Two smaller details are worth recording, with their limits. The grass is described as something “which today exists, and tomorrow is cast into the oven”, so the passage assumes its audience knew the oven as a place where plant matter was burned. The question “Which of you by being anxious can add a cubit to his height?” uses a unit of length without defining it. The text gives neither the cubit’s size nor the oven’s type.[1]

The passage ends on a claim rather than an artefact: “your Father knows that you need these things.” That is the text’s own theological statement.[1] ■

SOURCES 1 Luke 12:22-31 (WEB) bkmk.to/3f1

Sowing, reaping, barns: The Harvesters, 1565

PIETER BRUEGEL THE ELDER, THE HARVESTERS (1565). CC0 1.0, THE METROPOLITAN MUSEUM OF ART (OPEN ACCESS)

“That tells us how the passage pictures security, namely a full barn.”

WORTH YOUR TIME

Two long reads

SCAN THE CODE OR TYPE THE SHORT LINK. ALL BOOKMARKED THIS WEEK; TOGETHER ABOUT 6 MINUTES

1 4 MIN · REPORTAGE @TheHackersNews THE HACKER NEWS

New QPACK Flaws Can Crash HTTP/3 Servers, Trigger Up to 4,000-Fold Memory Amplification

A compact technical report on a flaw FoxIO found in seven of the 10 HTTP/3 stacks it tested.

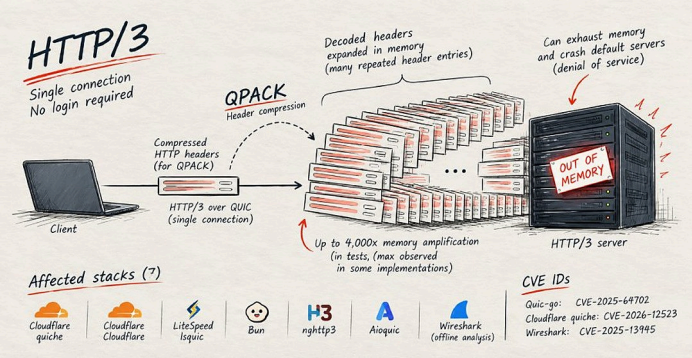

GRAPHIC: @THEHACKERSNEWS

It explains the QPACK header-compression cause, gives measured amplification figures and includes a table of fixed and unpatched stacks. Bun, the h3 Rust library and aioquic were still unpatched at disclosure.

bkmk.to/3r1

"The attacker needs no account or special access, and since HTTP/3 runs over UDP, the request reaches the server before much of the usual filtering applies."

THE HACKER NEWS · *New QPACK Flaws Can Crash HTTP/3 Servers, Trigger Up to 4,000-Fold Memory Amplification*

2

Pope discharged from hospital days after lung cancer surgery

A brisk account of Pope Leo XIV leaving Rome's Gemelli Hospital four days after surgery to remove the lower lobe of his right lung, and notes that results of tests for any spread of the cancer have not been announced. It also reports his call with Vladimir Putin, the first since June 2025, and a Latin America trip planned for early November. The BBC says it is not yet clear whether he will be fit to travel.

2 MIN · REPORTAGE @BBCBreaking BBC NEWS

bkmk.to/3r2

ELSEWHERE THIS ISSUE

3 Write the memo yourself

4 SpaceX's Super Intelligence arm now outsells its space business

5 Anthropic outgrew git, so it built Scion

6 A near-record index with 206 battered stocks

ALSO BOOKMARKED SHORTER, STILL GOOD

10 fresh everyday fragrances for men
@manly_talks
bkmk.to/3k25

11/10 classic recommendations
@Humanoholic
bkmk.to/3k26

A man quits caffeine, saves his career
@ThorTorrens
bkmk.to/3k27

G.K. Chesterton on mysticism
@byistra
bkmk.to/3k28

Horse Trading
@ALG2Horse

POST

- ich, 34, Sozialkompetenz einer Ravioli-Dose - erster Elternabend 6b, Frau hat Spätschicht, ich muss allein - Einladung per Zettel im Ranzen, zerknüllt unter einer Banane von letzter Woche - "Wir freuen uns auf einen regen Austausch :)" - der Smiley. Eine Kriegserklärung - nehm den Zettel mit, falls ich nachlesen muss wo der Raum is - Plan: hinten sitzen, nicken, nix sagen, um 20:15 wieder im Auto - ein Meisterplan ohne Schwachstellen - komm rein, alle Tische im Kreis - es gibt kein…

16 OCT 2026 bkmk.to/3z1

POST: @ALG2HORSE

ALSO THIS WEEK

A Polish commentator's verdict on France

Krzysztof Szczawinski (@Kristof_Poland) wrote to French readers that what they are watching is not a crisis but the result of at least 45 years of bad decisions, which he blamed on bureaucratic socialism. Quoting a post about Marseille, he urged France to hand power back to builders, citing Poland's own experience.

@Kristof_Poland bkmk.to/3k22

ELTERNABEND

The parents' evening has no back row

A post by @ALG2Horse about a father's first parents' evening alone turns into a comedy of a child's chair, a costed school trip and an opposition nobody knew existed.

PROMPTED BY @ALG2Horse

A parents' evening invitation can arrive under a banana. On 16 October, @ALG2Horse (Horse Trading) wrote that the invitation to his first Elternabend, for class 6b, came home in the school bag "zerknüllt unter einer Banane von letzter Woche". His wife was on a late shift, so he went alone. He rated himself, at 34, as having the "Sozialkompetenz einer Ravioli-Dose".[1]

He arrived with a plan: "hinten sitzen, nicken, nix sagen, um 20:15 wieder im Auto". The tables, he wrote, stood in a circle, and "es gibt kein hinten". The only free chair was a primary-school one, seat height 34cm, which left his knees at chin height. Opposite sat six mothers with thermos mugs and cardigans, and a shared WhatsApp group with 400 unread messages he knew nothing about. Everyone else had the invitation neatly in a plastic sleeve.[1] His smelled faintly fermented.[1] He read the invitation's smiley as "Eine Kriegserklärung".[1]

> *"Wir waren die ganze Zeit eine Opposition, wir wussten es nur nicht"*
>
> @ALG2HORSE

The trouble began with the money. A class trip to Thüringen was to cost €389.50 per child. One mother wanted organic catering for €47 more, another a climbing-park afternoon for €32, and a third a mindfulness seminar at €28.90. Everyone nodded. His head did the sums by reflex, and out it came, too loudly: "Das sind dann 497,40€ für fünf Tage Thüringen, dafür flieg ich mit der ganzen Familie nach Malle". The arithmetic holds. Six thermos mugs stopped in mid-air. Leaning back to look casual, he forgot the chair, tipped over and landed on his back, legs over the table, socks holed at both big toes.[1]

Then a father at the back, whom he had not noticed because he sat just as low, said: "Der Mann hat recht". Another agreed. There were three fathers in all; the quiet minority had been there from the start. In the vote, the organic food and the mindfulness seminar went, while the climbing park stayed, because the children would otherwise mutiny. The total fell to €421.50, a saving of €75.90 per child across 27 children, or €2,049.30.[1]

The reward for numeracy was work. A pair of mothers approached him afterwards: "Wir suchen noch jemanden für die Klassenkasse, Sie können ja offensichtlich rechnen". He said yes, he wrote, because he was still on the floor and they were asking from above. He is now in the WhatsApp group, with 412 unread messages. His son reports that the other children call him the Malle-Papa. The teacher's minutes list his hobby as eating and his role as treasurer.[1]

The post is played for laughs, but its sharpest point concerns consensus. A circle of nodding heads can look like agreement while concealing dissent. The back row he hoped to hide in did not exist; the allies did. Saying the number aloud cost him a chair and his dignity, and it won him the class kitty.[1]

SOURCES 1 @ALG2Horse bkmk.to/3z1

The Weekend Bookmark

A personal weekly digest of posts bookmarked on X. Every story credits the accounts it came from; numbered notes give the short links.

THIS ISSUE
No. 3, Saturday, 24 October 2026

NEXT ISSUE
Saturday, 31 October 2026

TYPE
Newsreader, Source Serif 4, IBM Plex Sans Condensed and IBM Plex Mono.

PICTURES
p. 1: Illustration: vector drawing generated for this issue; p. 3: Illustration: vector drawing generated for this issue; p. 5: @doooyle on X; p. 9: @Techaktien1 on X; every other picture is credited beside it.

PRINTING
Sixteen A4 pages, double-sided, flip on long edge. Booklet: four A3 sheets, flip on short edge, fold and staple. Recycle when read.